



AI PROMPTING MASTERY

Trainer Session Kit

A one-to-one enablement session for a public-speaking coach building an AI-powered career-assessment tool for IIT graduates and senior professionals.

PREPARED FOR

HAVOX INTERNATIONAL

Public Speaking Trainer — IIT & Senior Professional Cohorts Session Design · Prompt Frameworks · Handout

SESSION OVERVIEW

Why This Session, and What He Actually Needs

He already uses ChatGPT and Gemini for content, so this is not an AI-literacy session. The real goal is to move him from casual prompting to **structured, repeatable prompt engineering** — specifically for two outputs: (1) his own training content for IIT-grad and senior-professional audiences, and (2) the logic behind a personalized career-assessment site, where a questionnaire feeds an AI model that generates a tailored career report.

Session Objectives

- He can build a prompt from a repeatable framework instead of ad-hoc typing.
- He knows which AI model to reach for depending on the task (reasoning vs. speed vs. tone control).
- He understands why questionnaire design determines report quality, not the prompt alone.
- He can write a working master prompt that turns structured answers into a personalized report.
- He knows how to catch hallucinated or generic output before a user ever sees it.
- He can frame career suggestions responsibly, without overpromising or exposing user data carelessly.
- He can reuse the same framework for his own content (LinkedIn, emails, recaps) — not just the tool.
- He leaves with a one-page brief format he can hand to a developer to build an MVP, not a full product.

Full Content Map for Today

- 1. Prompt frameworks & techniques (R-C-T-O-N-E, few-shot, chain-of-thought, etc.)
- 2. AI model literacy — ChatGPT vs. Gemini vs. Claude, and when to use which
- 3. Questionnaire & data architecture — question design, scoring logic
- 4. Output reliability & hallucination control — closed-set grounding, edge-case testing
- 5. Personalization vs. generic output — killing "horoscope energy"
- 6. Ethics & responsible framing — disclaimers, overpromising, data privacy
- 7. Turning this into a repeatable content engine for his own work
- 8. Handoff & non-technical product thinking — briefing a developer, MVP scope

Session Flow (2 hours)

Time	Block	What Happens
0–10 min	Diagnose	Review his current prompts on ChatGPT/Gemini. Identify what's missing: no role, no format, no constraints.
10–30 min	Prompt Framework	Teach the R-C-T-O-N-E prompt framework. Build one prompt live, together, for his training content.
30–40 min	Model Literacy	Walk through ChatGPT vs. Gemini vs. Claude — strengths, and which to use for the career-report engine.
40–60 min	Career-Tool Logic	Map the questionnaire → scoring → prompt → report pipeline. Draft 10–12 questionnaire items together.
60–75 min	Master Prompt	Write the master AI prompt that ingests questionnaire answers and outputs a structured career report.
75–85 min	Prompt Techniques	Walk through chain-of-thought, few-shot, and other techniques with his own examples.

85–95 min	Reliability & Personalization	Stress-test the prompt with edge-case answers. Fix generic/hallucinated output using grounding and specificity.
95–105 min	Ethics & Framing	Set disclaimers, tone limits, and a simple data-privacy stance for the tool.
105–115 min	Content Engine & Handoff	Show how the same framework powers his LinkedIn/emails. Build the developer brief for an MVP.
115–120 min	Wrap & Handout	Hand him the one-pager: framework + prompt template + developer brief format.

What To Watch For While Teaching Him

- Is he writing prompts as one long paragraph? → Push him to separate Role / Context / Task / Format.
- Is he accepting the first AI output as final? → Force one round of critique-and-refine on every example.
- Is he designing questions that only he can interpret (not AI)? → Questions must map to explicit traits/tags.
- Is he thinking about tone-matching (sounding like a career report, not a chatbot)? → Show output-format constraints.
- Does he understand negative prompting (what to exclude)? → This is usually the most skipped skill — spend real time here.
- Can he explain back the framework in his own words before you move on? → Don't advance until he can.

PROBLEM FRAMING

PRD-Style Brief: The Career-Assessment Tool

Before touching prompts, get him to define the problem like a product brief. This is what makes his AI output consistent instead of random. Walk through this table with him live.

Field	Definition For This Tool
Problem	Professionals don't know which field/role fits them; generic career advice is not personalized.
User	IIT graduates and senior professionals attending his public-speaking / career sessions.
Input	A structured questionnaire (interests, strengths, work-style, values, risk appetite, past wins).
Process	Answers are tagged/scored → formatted into a data block → inserted into a master AI prompt.
Output	A personalized report: suggested fields, strengths, growth areas, next steps, sample career paths.
Constraints	Must sound professional and specific, not generic horoscope-style text. Must fit on one web page.
Success	User reads the report and feels 'this is clearly about me,' not 'this could be about anyone.'

His Role vs. The Developer's Role

- **He owns:** the questionnaire design, the scoring logic, the tone/voice of reports, and the master prompt.
- **Developer owns:** the website, the form UI, sending answers to the AI API, and displaying the output.
- The output of this session is the part he owns — everything AI-related — packaged so a developer can wire it up.

Deliverable Framework — What He Walks Away With

#	Deliverable	Format
1	Prompt framework (reusable for any content he creates)	1-page reference card
2	Questionnaire draft (10–12 questions, tagged to traits)	Table / list
3	Master prompt template for the career report	Copy-paste prompt block
4	Sample generated report (run live in session)	Text output
5	Developer brief (what to hand off to build the site)	1-page brief

CORE FRAMEWORK

The R-C-T-O-N-E Prompt Framework

Teach this as the single framework he uses for everything — training content and the career-report prompt alike. Six components, always in this order.

R — ROLE

Who the AI should act as. Sets vocabulary, authority, and perspective.

"You are a senior career counselor who has advised 500+ engineering graduates from IITs on career transitions."

C — CONTEXT

Background the AI needs but shouldn't guess. Audience, situation, prior data.

"The user is a 24-year-old mechanical engineering graduate who just completed a 12-question self-assessment."

T — TASK / CONTENT

The exact action, and the raw input data it should work on.

"Analyze the following questionnaire answers and identify the 3 career fields that best match this person: [DATA]"

O — OUTPUT FORMAT

Structure, length, headings — so it's usable, not a wall of text.

"Return the report in this exact structure: Top 3 Fields / Why It Fits You / Strengths / Growth Areas / Next Step. Max 350 words."

N — NEGATIVE PROMPTING

What to exclude. This is the most skipped step — spend real time here.

"Do not use generic horoscope language. Do not suggest more than 3 fields. Do not use the words 'passionate' or 'synergy'."

E — EXAMPLES

1–2 sample outputs so the AI matches tone and depth, not just structure.

"Here is a sample report we consider a 5/5 in tone and depth: [PASTE EXAMPLE REPORT]"

WORKED EXAMPLE

Full Framework Applied — Career Report Master Prompt

This is the live example to build together in the session. Walk through each line and ask him to identify which of the 6 components it belongs to before you tell him.

You are a senior career counselor and industrial psychologist who has advised over 500 engineering and management graduates from IITs and top Indian institutes on career direction.

CONTEXT:

The user has just completed a 12-question self-assessment covering interests, work-style, risk appetite, values, and past achievements. Their raw answers and derived trait scores are provided below.

TASK:

Analyze the trait scores and answers. Identify the 3 career fields that best match this person, ranked by fit. For each field, explain WHY it fits, using specific evidence from their answers (not generic statements).

USER DATA:

```
{{questionnaire_answers}}
{{trait_scores}}
```

OUTPUT FORMAT (follow exactly):

1. Top 3 Fields (ranked, one line each)
2. Why It Fits You (2-3 sentences per field, referencing their actual answers)
3. Your Core Strengths (3 bullet points)
4. One Growth Area to Watch
5. Recommended Next Step (one concrete action)

Total length: 300-350 words. Second-person voice ("you"), professional tone.

DO NOT:

- Use generic/horoscope-style language ("you are a natural leader")
- Recommend more than 3 fields
- Use filler words: "passionate," "synergy," "dynamic," "go-getter"
- Repeat the questionnaire questions back to the user
- Mention that you are an AI or that this was generated by a prompt

EXAMPLE OF DESIRED TONE (for reference only, do not copy content):

```
{{paste_one_sample_report_here}}
```

Point out to him: the `{{ }}` placeholders are exactly what the website's form will inject automatically once a developer builds it — this is the bridge between his prompt and the actual product.

QUESTIONNAIRE DESIGN

Mapping Questions to Traits (Build This Together)

The AI report is only as good as the data it receives. Draft questions that map to explicit, nameable traits — not vague ones. Use this table as a live worksheet.

Trait Tag	Sample Question	What It Reveals
Analytical vs. Creative	When solving a new problem, do you prefer a proven method or inventing your own approach?	Field fit: engineering/finance vs. design/strategy
People vs. Systems	Do you get more energy from a great conversation or from fixing a broken process?	People-facing roles vs. technical/ops roles
Risk Appetite	Would you rather a stable, well-paid role or a high-upside, uncertain one?	Startup/entrepreneurial fit vs. corporate fit
Depth vs. Breadth	Do you enjoy mastering one skill deeply, or juggling many different skills?	Specialist vs. generalist career paths
Past Wins	Describe one achievement you're proud of and why it mattered to you.	Real evidence the AI can quote back specifically
Values	What matters more: impact, income, recognition, or autonomy? Rank them.	Filters which fields will feel meaningful, not just suitable

Have him extend this to 10–12 questions. Rule to enforce: **every question must map to a trait tag that gets fed into the prompt as data** — no question should exist just because it "sounds nice."

PROMPT TECHNIQUES

Prompting Techniques He Should Know

After the framework, layer in these techniques. He doesn't need all of them for every task — teach him to recognize which one fits which problem.

Zero-Shot Prompting

Ask directly with no examples. Fastest, works for simple/familiar tasks.

"Summarize this speech feedback in 3 bullet points."

Few-Shot Prompting

Give 2–3 examples of input→output so AI matches his exact style and tone.

"Here are 2 example career reports I've written. Write a third, same tone, for this new data: [DATA]"

Chain-of-Thought (CoT)

Ask AI to reason step-by-step before giving the final answer. Improves accuracy on judgment calls like ranking career fields.

"Think step-by-step: first identify the top 2 traits, then match each to a field, then rank by fit. Show your reasoning, then give the final answer."

Role Prompting

Assign a specific persona to shift vocabulary and depth (covered in the R of the framework, but worth naming separately).

"Act as a skeptical hiring manager reviewing this candidate's profile."

Self-Critique / Refinement

Ask the AI to review and improve its own output — catches generic or repetitive language.

"Review your last answer. Rewrite it removing any generic phrases, and make it feel written specifically for this person."

Constrained / Template Prompting

Force strict output structure (JSON, fixed headings) so it can be dropped into a website automatically.

"Return your answer ONLY as valid JSON with keys: top_fields, strengths, growth_area, next_step."

Iterative Refinement

Never accept the first draft. Re-prompt with specific fixes rather than starting over.

"Good, but point 2 is too generic – make it reference their specific answer about teamwork."

MODEL LITERACY

Which AI Model to Use, and When

He already uses ChatGPT and Gemini. Show him this isn't about brand loyalty — different models behave differently on reasoning depth, tone control, and honesty, and the career-report tool should be built on whichever handles structured, personalized output most reliably.

Model	Strength	Best Use Here
ChatGPT (GPT)	Fast, strong at conversational and creative content.	Drafting social posts, quick content, brainstorming.
Gemini	Strong at summarizing and working with structured/long input.	Digesting long questionnaire data before drafting a report.
Claude	Strong instruction-following, better at sticking to strict formats and avoiding generic filler.	The master career-report prompt — where format and negative prompting must hold exactly.

Exercise: run the exact same master prompt on two different models with identical questionnaire data, and compare which one avoids generic language better. This teaches him to evaluate output, not just generate it.

Reasoning Mode vs. Fast Mode

- Fast/default mode: fine for short content — LinkedIn captions, email drafts.
- Reasoning/"thinking" mode: use for the career report — ranking 3 fields out of many is a judgment call, and reasoning models catch contradictions in the user's answers better.

OUTPUT RELIABILITY

Catching Hallucinated or Generic Output

AI will confidently suggest a career field that doesn't actually fit, or default to safe, generic phrasing. This section teaches him to build in guardrails rather than trust output blindly.

Grounding: Closed-Set vs. Open Generation

- **Open generation** (risky): "Suggest career fields for this person." AI can invent vague or irrelevant fields.
- **Closed-set grounding** (reliable): "Choose the top 3 fields ONLY from this list: [Engineering, Product Management, Consulting, Data Science, Design, Sales, Finance, Academia, Entrepreneurship...]" — this keeps output relevant and comparable across users.

Edge-Case Testing (Do This Live in Session)

- Run the prompt with a **contradictory** answer set (e.g., says they love people AND hate teamwork) — does the report still make sense?
- Run it with a **minimal/lazy** answer set (one-word answers) — does AI admit low confidence, or does it fabricate depth it doesn't have?
- Run it twice with the exact same input — is the output stable, or wildly different each time? Instability signals the prompt needs tighter constraints.

RULE OF THUMB

If the AI can't justify a suggestion by pointing to a specific answer the user gave, the suggestion shouldn't be in the report. Add this as an explicit instruction in the prompt itself.

PERSONALIZATION

Killing "Horoscope Energy"

The single biggest quality risk in this tool: reports that sound personalized but could apply to almost anyone. Teach him to test for this specifically.

The Swap Test

Take a finished report and try mentally swapping in a different user's name and answers. If the report still reads fine, it's too generic — it isn't actually using their specific data.

Techniques to Force Specificity

- Instruct the AI to quote or closely reference the user's own words from open-ended answers.
- Require at least one sentence per section that could only be true for this specific person.
- Ban vague trait words unless tied to evidence: not "you are a natural leader" but "your answer about resolving the team conflict in Q7 shows leadership under pressure."
- Use self-critique (from the techniques page): "Rewrite this removing any sentence that could apply to a different person."

ETHICS & FRAMING

Responsible Career Advice, Not Prophecy

This tool is advising real people on career direction. Get him to bake in responsible framing from day one, not as an afterthought.

Language Discipline

- Say **"aligns with your traits"**, not **"you will succeed in X."** Directional, not deterministic.
- Every report should carry a short disclaimer: this is a starting point for reflection, not a verdict.
- Avoid absolute language: "never," "guaranteed," "perfect fit."

Data Privacy — Ask Him These Questions Directly

- What happens to a user's questionnaire answers after the report is generated — stored, deleted, or anonymized?
- Will answers ever be reused to train or improve anything, and does the user know that?
- Is there a clear line the user can read before submitting, saying what will and won't be done with their data?

WHY THIS MATTERS FOR HIM SPECIFICALLY

His users are senior professionals and IIT graduates — a credibility-sensitive audience. One overconfident or careless report undermines trust in the entire tool and his personal brand.

BEYOND THE TOOL

A Repeatable Content Engine — and the MVP Handoff

Close the session by showing him this framework isn't just for the career tool — it's now his default way of using AI for everything else he creates.

Reusing the Framework for His Own Content

- LinkedIn posts recapping a training session — same R-C-T-O-N-E structure, different Role/Output.
- Follow-up emails to workshop attendees — Few-shot using his past emails as examples.
- Session recaps and internal notes — Chain-of-thought to structure key takeaways logically.
- Build a running "prompt library" doc: every time a prompt works well, he saves it as a template for reuse.

Briefing a Developer for an MVP (Not the Full Vision)

- MVP = a simple form (Google Form or Typeform) + his master prompt run manually or via a no-code tool (e.g., Glide, Bubble, or a basic API wrapper) — not a fully custom-built platform on day one.
- He should hand off exactly: the finalized questionnaire, the scoring/tagging logic, the master prompt with placeholders, 2–3 approved sample reports, and the required output format.
- Test the MVP with 5–10 real users before investing in a polished website — this avoids building the wrong thing well.

CLOSING POINT FOR THE SESSION

He doesn't need to become technical. He needs to own the judgment calls — what to ask, what tone to demand, what to forbid, and how to verify the output is honest. That's the entire job; everything else can be outsourced.

TAKEAWAY HANDOUT

One-Page Handout — Give This to Him at the End

Print or send this page separately as his leave-behind reference. It condenses the whole session into what he needs at his desk.

1. The Framework (Every Prompt, Every Time)

- Role — who is the AI acting as?
- Context — what background does it need?
- Task/Content — what exactly should it do, with what data?
- Output format — what structure, length, tone?
- Negative prompting — what must it avoid?
- Examples — what does "good" look like?

2. Techniques Quick-Reference

Technique	Use When...
Zero-shot	The task is simple and familiar
Few-shot	You need it to match your exact tone/style
Chain-of-thought	The task needs ranking, judgment, or comparison
Self-critique	The first draft feels generic
Constrained/template	Output needs to plug into a website or app
Iterative refinement	Always — never accept draft one

3. Quality & Ethics Checklist (Run Before Publishing Any Report Template)

- Grounded to a closed list of fields — not open-ended guessing.
- Passes the Swap Test — could not apply to a different user unchanged.
- Tested against a contradictory and a minimal answer set.
- Uses "aligns with" language, not "you will succeed" language.
- States what happens to the user's data, in plain language, before they submit.

4. Developer Hand-Off Checklist (MVP, Not Full Product)

- Final questionnaire (10–12 questions, each tagged to a trait).
- Scoring/tagging logic — how answers convert into trait data.
- Master prompt template, with {{placeholders}} marked clearly.
- 2–3 sample reports approved by him as "this is the tone I want."
- Required output format (plain text vs. JSON) for the developer to display it correctly.

REMEMBER

The AI is only as good as (1) the questions you ask the user, and (2) the constraints you give the AI. Most bad AI output comes from skipping Negative Prompting and Examples — don't skip them.

Havox International · AI Prompting Mastery Trainer Session Kit